

# Inter-Processor Communication Design

## IPC Concepts



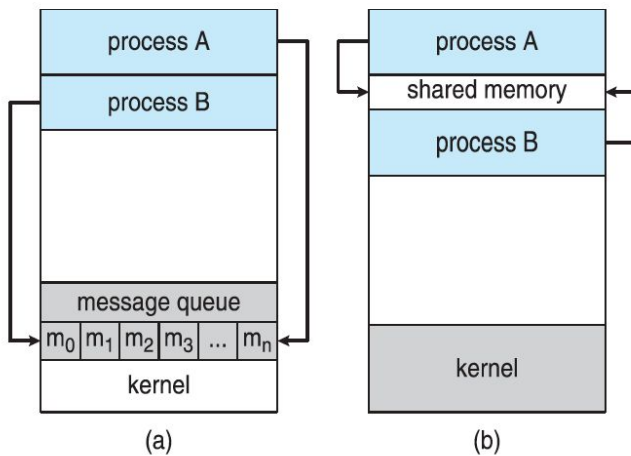
### Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
  - **Shared memory**
  - **Message passing**



### Communications Models

(a) Message passing. (b) shared memory.





## Producer-Consumer Problem

---

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - **unbounded-buffer** places no practical limit on the size of the buffer
  - **bounded-buffer** assumes that there is a fixed buffer size



## Interprocess Communication – Shared Memory

---

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.



## Interprocess Communication – Message Passing

---

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send(message)**
  - **receive(message)**
- The *message* size is either fixed or variable



## Message Passing (Cont.)

---

- If processes *P* and *Q* wish to communicate, they need to:
  - Establish a **communication link** between them
  - Exchange messages via send/receive
- Implementation issues:
  - How are links established?
  - Can a link be associated with more than two processes?
  - How many links can there be between every pair of communicating processes?
  - What is the capacity of a link?
  - Is the size of a message that the link can accommodate fixed or variable?
  - Is a link unidirectional or bi-directional?



## Direct Communication

---

- Processes must name each other explicitly:
  - `send(P, message)` – send a message to process P
  - `receive(Q, message)` – receive a message from process Q
- Properties of communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional



## Indirect Communication

---

- Mailbox sharing
  - $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
  - $P_1$  sends;  $P_2$  and  $P_3$  receive
  - Who gets the message?
- Solutions
  - Allow a link to be associated with at most two processes
  - Allow only one process at a time to execute a receive operation
  - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



## Synchronization

---

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
  - **Blocking send** -- the sender is blocked until the message is received
  - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
  - **Non-blocking send** -- the sender sends the message and continue
  - **Non-blocking receive** -- the receiver receives:
    - A valid message, or
    - Null message
- Different combinations possible
  - If both send and receive are blocking, we have a **rendezvous**



# Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
  1. Zero capacity – no messages are queued on a link.  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits

## Message Passing Models:

From a popular message queue tool: RabbitMQ

<https://www.rabbitmq.com/getstarted.html>

### 1 "Hello World!"

The simplest thing that does  
*something*




Python

Java


### 2 Work queues

Distributing tasks among  
workers (the competing  
consumers pattern)



### 3 Publish/Subscribe


Sending messages to many  
consumers at once



Python


### 4 Routing

Receiving messages  
selectively




### 5 Topics

Receiving messages based on  
a pattern (topics)



### 6 RPC

Request/reply pattern  
example



## IPC in Pi-OS: Topic-based message passing queue

The pi-OS provide the ability for applications to communicate. The kernel will keep a queue of topics, each topic may have several subscribers. Each task can subscribe some topics. Each task has a message queue to store all the messages of booked topics. When a sender send a message of topic\_A, the kernel is responsible for passing this message to all the subscribers.



Core data structures:

```
64 #define TOPIC_QUEUE_LEN 8
65 static topic_struct * topic_struct_list[TOPIC_QUEUE_LEN];

13 typedef struct Tmessage_struct{
14     char topic[32];
15     char sender_name[32];
16     int points; // Receiver can get points from this message
17     int msg_id;
18     char msg[64];
19 }message_struct;
20
21 #define MAX_SUBSCRIBER_NUM 10
22 typedef struct Ttopic_struct{
23     char topic[32];
24     int subscriber_number;
25     struct Ttask_struct* owner;
26     struct Ttask_struct* subscriber_list[MAX_SUBSCRIBER_NUM];
27 }topic_struct;

56 #define TASK_MSG_QUEUE_LEN 3
57
58 typedef struct Ttask_struct{
59     int state;
60     int prio;
61     int policy;
62     struct Ttask_struct* parent;
63     struct list_head tasks;
64     int pid;
65     void (*RunFunc)();
66     void (*tasks_entry)(struct Ttask_struct * pTask);
67     unsigned WakeTicks;
68     TTaskRegisters Regs;
69     u8 *pStack;
70     unsigned StackSize;
71     int active_child_count;
72     int isWaiting;
73     message_struct msgQueue[TASK_MSG_QUEUE_LEN];
74 }task_struct;
```

## Important Functions

A task can create topic, subscribe a topic, send message to a topic group, and receive message.

**Create topic:** kernel will create a new topic struct, and put it in the topic\_struct\_list. This topic struct will be initialized to have zero subscribers.

**Subscribe topic:** kernel will put requesting task in the subscriber\_list of a certain topic.

**Send(topic, message\*):** kernel will copy this message to all the subscribers in this topic subscriber list. So this copy will be put in each subscribing task message queue. If queue full, oldest message will be replaced.

**Receive(topic, message\*):** Task want to receive message from its message queue. Kernel will choose the oldest message in the queue, copy it. And remove this message in the message queue.

## Mission:

Implement ipcSend(void \* msg\_send) & ipcReceive(void \* msg\_receive)

```
267 //Copy this message to all the subscribers' message queue;
268 //If msg queue full, replace the oldest one
269 int CKernel::ipcSend(void * msg_send){
270     //write_log("now in CKernel::ipcSend ");
271     message_struct *pMsg = (message_struct *) msg_send;
272     //TODO
273
274     return 0;
275 }
276
277 //Non-blocking receive; read message from message queue to msg_receive,
278 //read the oldest one and clear that space in message queue
279 int CKernel::ipcReceive(void * msg_receive){
280     //write_log("now in CKernel::ipcReceive ");
281     message_struct *pMsg = (message_struct *) msg_receive;
282     task_struct * pTask = getScheduler()->pCurrentTask;
283     //TODO
284
285     return 0;
286 }
---
```

The codes you need to add is about 50 lines, so it's mainly understanding work.

## Tasks

There are four user tasks in project 6, professor task, student01 task, student02 task, student03 task, and their behavior is defined:

```

87 void task_professor_run()
88 {
89     message_struct msg;
90     printf("PROFESSOR: Now we are in task_professor_run *****");
91
92     ipc_create(TOPIC_01);
93     ipc_create(TOPIC_02);
94
95     sleep(2);
96     printf("PROFESSOR: lecturing");
97     memset(&msg, 0, sizeof(msg));
98     strcpy(msg.topic, TOPIC_01);
99     strcpy(msg.msg, HOMEWORK_01);
100    msg.points = 20;
101    ipc_send(&msg);
102    sleep(1);
103    memset(&msg, 0, sizeof(msg));
104    strcpy(msg.topic, TOPIC_02);
105    strcpy(msg.msg, PROJECT_01);
106    msg.points = 30;
107    ipc_send(&msg);
108    sleep(1);
109    printf("PROFESSOR: lecturing");
110    memset(&msg, 0, sizeof(msg));
111    strcpy(msg.topic, TOPIC_01);
112    strcpy(msg.msg, HOMEWORK_02);
113    msg.points = 20;
114    ipc_send(&msg);
115
116 void task_student01_run()
117 {
118     printf("STUDENT_01: Now we are in task_student01_run *****");
119     sleep(1);
120     ipc_subscribe(TOPIC_01);
121     ipc_subscribe(TOPIC_02);
122     sleep(2);
123     message_struct msg;
124     char log[200];
125     memset(log, 0, 200);
126     void * pFile = open("stu01.log", 1);
127     seek(pFile, 0);
128     int recv_flag = 1;
129     int course_to_wait = 2;
130     int grade = 0;
131     while(course_to_wait){
132         memset(&msg, 0, sizeof(message_struct));
133         recv_flag = ipc_receive(&msg);
134         if(1==recv_flag){
135             printf("Student01 receive no msg, sleep ");
136             sleep(2);
137             continue;
138         }
139         if(0==msg.points){
140             printf("Student01 receive msg of %s, content: %s, get points %d, msg=%d", msg.topic, msg.msg,
141                 course_to_wait --, msg.points);
142             write(pFile, msg.topic, strlen(msg.topic));
143             write(pFile, " ", 2);
144             write(pFile, msg.msg, strlen(msg.msg));
145             write(pFile, "\n", 1);
146             continue;
147         }
148     }
149     printf("Student01 receive msg of %s, content: %s, get points %d", msg.topic, msg.msg, msg.points);
150     write(pFile, msg.topic, strlen(msg.topic));

```



## Sample output:

```
00:00:07.29 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:08.20 kernel: SCHEDULER:: yield to task:professor, Index=0, PID=101, PPID=101, state=0---
00:00:08.21 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:08.25 kernel: SCHEDULER:: yield to task:student 01, Index=1, PID=102, PPID=102, state=0---
00:00:08.26 kernel: Student01 receive msg of ENEE447, content: Project 3, get points 40
00:00:08.27 kernel: SCHEDULER:: yield to task:student 02, Index=2, PID=103, PPID=103, state=0---
00:00:08.28 kernel: Student02 receive msg of ENEE350, content: Homework 04, get points 20
00:00:08.29 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:09.21 kernel: SCHEDULER:: yield to task:professor, Index=0, PID=101, PPID=101, state=0---
00:00:09.22 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:09.27 kernel: SCHEDULER:: yield to task:student 01, Index=1, PID=102, PPID=102, state=0---
00:00:09.28 kernel: Student01 receive msg of ENEE350, content: Homework 05, get points 20
00:00:09.29 kernel: SCHEDULER:: yield to task:student 03, Index=3, PID=104, PPID=104, state=0---
00:00:09.30 kernel: Student03 receive msg of ENEE447, content: Course ends, get points 0,
00:00:09.31 kernel: Student03 grade is 100
00:00:09.32 kernel: log is:
ENE447 Project 1
ENE447 Project 2
ENE447 Project 3
ENE447 Course ends

00:00:09.34 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:10.22 kernel: SCHEDULER:: yield to task:professor, Index=0, PID=101, PPID=101, state=0---
00:00:10.23 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:10.29 kernel: SCHEDULER:: yield to task:student 01, Index=1, PID=102, PPID=102, state=0---
00:00:10.30 kernel: Student01 receive msg of ENEE350, content: Course ends, get points 0,
00:00:10.31 kernel: Student01 receive msg of ENEE447, content: Course ends, get points 0,
00:00:10.32 kernel: Student01 grade is 150
00:00:10.33 kernel: log is:
ENE350 Homework 01
ENE350 Homework 02
ENE447 Project 2
ENE350 Homework 04
ENE447 Project 3
ENE350 Homework 05
ENE350 Course ends
ENE447 Course ends

00:00:10.35 kernel: SCHEDULER:: yield to task:student 02, Index=2, PID=103, PPID=103, state=0---
00:00:10.36 kernel: Student02 receive msg of ENEE350, content: Homework 05, get points 20
00:00:10.37 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
00:00:12.37 kernel: SCHEDULER:: yield to task:student 02, Index=2, PID=103, PPID=103, state=0---
00:00:12.38 kernel: Student02 receive msg of ENEE350, content: Course ends, get points 0,
00:00:12.39 kernel: Student02 grade is 80
00:00:12.40 kernel: log is:
ENE350 Homework 01
ENE350 Homework 02
ENE350 Homework 04
ENE350 Homework 05
ENE350 Course ends

00:00:12.42 kernel: SCHEDULER:: no READY task in queue now, scheduler quite ----
```

DELL



## Goal:

Understanding IPC and message queue

Implementat send & receive